

The Module Language

Part I – Modules, Compilation, Applying Functors

OCaml **PRO** – Benjamin Canou

March 9-13 2015

Outline

Toplevel Modules
Abstraction
The module Language
Using Standard Functors

Toplevel Modules

- Source units as modules
- Separate compilation
- Interfaces
- Interface compilation
- Documentation

Source units as modules

Basics:

- Each file `momo.ml` defines a module named `Momo`.
- Its names can be accessed from other modules prefixed by `Momo`.
- The module relationship is a partial order (cycles are forbidden).

File `momo.ml`:

```
1 : type t = { x : int ; y : int }  
2 : let origin = { x = 0 ; y = 0 }  
3 : let above pa pb = pa.x >= pb.x
```

File `main.ml`:

```
1 : let o : Momo.t = Momo.origin  
2 : let p : Momo.t = { Momo.x = 2 ; Momo.y = 4 }  
3 : let () = assert (Momo.above p o)
```

Prefixes can be omitted by (very cautiously) using

- `open Momo` at toplevel
- `let open Momo in` in expressions
- `Momo.(...)` in expressions

Separate compilation

Multi-modular programs can be simply compiled in one invocation:

```
ocamlopt momo.ml main.ml -o momotest
```

Order is important !

Or modules can be compiled separately, and linked together:

```
ocamlopt -c momo.ml
```

```
ocamlopt -c main.ml
```

```
ocamlopt momo.cmx main.cmx -o momotest
```

A simple Makefile:

```
1 : momotest: momo.cmx main.cmx
2 :          ocamlopt $^ -o $@
3 : %.cmx: %.ml
4 :          ocamlopt -c $<
5 : clean:
6 :          -$(RM) -f *.cm* momotest
```

Interfaces

Each `.ml` source can be accompanied with a `.mli` interface.

Interfaces define:

- The types present in the source, with the same syntax.
- Each value (function or not) present in the source in the form:
`val name : type`

File `momo.mli`:

```
1: type t = { x : int ; y : int }  
2: val origin : t  
3: val above : t -> t -> bool
```

Interface compilation

Interfaces are compiled to `cmi` files.

- When checking a reference to an external module, `ocamlopt` will look for its `cmi` first.
- When an interface is present, `ocamlopt` will require it to be compiled before the source.

The `Makefile` updated:

```
1: momotest: momo.cmx main.cmx
2:          ocamlopt $^ -o $@
3: %.cmx: %.ml
4:          ocamlopt -c $<
5: %.cmx: %.mli
6:          ocamlopt -c $<
7: momo.cmx: momo.cmi
8: clean:
9:          -rm -f *.cm* momotest
```

Com more complex dependencies, there is `ocamldep`.

Documentation

The `ocamldoc` tool defines a special syntax for comments:

- Documentation comments are `(** enclosed like this *)`.
- Associated to the nearest element (before or after, without linebreak).
- The first of the file describes the module.

File `momo.mli` documented, `ocamldoc -html momo.mli` to see the result.

```
1 : (** The world of Momo *)
2 :
3 : (** The type of points in the world of Momo. *)
4 : type t =
5 :   { x : int (** How far on the horizon *) ;
6 :     y : int (** How high in the sky *) }
7 :
8 : (** The location of the mother of Momo. *)
9 : val origin : t
10 :
11 : (** To tell who is taller. *)
12 : val above : t -> t -> bool
```


Abstraction

Structure rewriting & hiding
Abstract types
Private types

The interface must not follow the implementation completely.

- To hide internal auxiliary elements.
- To hide functions that may break internal invariants.
- To help providing a more structured documentation.

For this, it is possible to:

- Hide types and values.
- Reorder everything.
- Use type aliases instead of primitive types.

Rewriting is allowed as long as

- The interface is still consistent with the implementation.
(An integer in the implementation cannot be exported with type float.)
- The interface is self consistent.
(types are used after their definition, etc.)

We speak of **structural inclusion** between:

- The internal structure, inferred by OCaml.
- The external structure, written by the programmer.

This inclusion is checked when compiling the implementation.

The definition of a type can be hidden, making it **abstract**.

- The syntax is `type t` (cut everything from the `=`).
- Abstract types cannot be constructed or destructed except in their own module.
- The structure is completely unknown of the outside.
- **Don't forget to provide builders and accessors !**

What you cannot do:

- Hide only some fields of a record.
- Hide only some constructors of a sum.

Why type abstraction ?

- Usual point: **to control the implementation.**
- To clarify the API documentation (not always the best solution).
- To preserve internal invariants.
- For program architecture: **write abstract interfaces first.**
- For quick prototyping (bosket oriented programming).

File `momo.mli`, its type `t` abstracted.

```
1 : type t
2 : val origin : t
3 : val above : t -> t -> bool
```

Note the regression: one cannot use pattern matching on `Momo.t`.

The previous `main.ml` is not valid anymore:

```
1 : let o : Momo.t = Momo.origin
2 : let p : Momo.t = { Momo.x = 2 ; Momo.y = 4 }
3 : let () = assert (Momo.above p o)
```

We have to rework `momo.mli`, to provide the necessary builders and accessors.

```
1 : type t
2 : val origin : t
3 : val translate : t -> int -> int -> t
4 : val above : t -> t -> bool
```

And update `main.ml`:

```
1 : let o : Momo.t = Momo.origin
2 : let p : Momo.t = Momo.(translate x 4 origin)
3 : let () = assert (Momo.above p o)
```

A less brutal alternative to type abstraction.

- Syntax: `type t = private ...`.
- The definition is kept public and appears in the API.
- Values cannot be constructed outside of the module.
- Values can be destructured outside of the module.

Differences with type abstraction:

- ⊖ Less possibilities of implementation rewrite.
- ⊖ Can preserve internal invariants.
- ⊕ Can make the documentation clearer.
- ⊕ Makes values destructible for the outside.
- ⊕ No need for accessors, just builders.

A same interface can contain a mix of public, abstract and private types.

File `momo.ml` remains unchanged:

```
1: type t = { x : int ; y : int }  
2: let origin = { x = 0 ; y = 0 }  
3: let translate { x ; y } dx dy = { x = x + dx ; y = y + dy }  
4: let above pa pb = pa.x >= pb.x
```

We add the `private` modifier to `momo.mli`.

```
1: type t = private { x : int ; y : int }  
2: val origin : t  
3: val translate : t -> int -> int -> t  
4: val above : t -> t -> bool
```

And update `main.ml`:

```
1: let o : Momo.t = Momo.origin  
2: let p : Momo.t = Momo.(translate x 4 origin)  
3: let () = assert (Momo.above p o)
```

The module Language

Local modules

Local signatures

Example: variations on a (key x value) table

Abstraction and signature rewriting

Composition

A module can define child modules:

- Syntax: `module Name = struct ... end.`
- They support the same features as toplevel modules, including modules.
- Access Syntax: `Module.Child.Child. ...`

That can be used for:

- Architecturing the API (by topic, hierarchically, etc.)
- Grouping local utilities to be hidden.
- Locally extending / patching external modules.
- Defining functor parameters (to be continued).

Local signatures

Child modules can also be restricted using signatures.

In the parent module signature:

- Syntax: `module Name : sig ... end.`
- Elements inside the `sig` are the same than in `mli` s.
- This way, the module is restricted only for the outside world.
- It is also possible to hide a child module completely.

Inside the module.

- A child module can also be restricted for the rest of the module.
- Syntax: `module Name : sig ... end = struct ... end`

Signatures can be named:

- Syntax: `module type NAME = sig ... end`
- For use in interfaces: `module Name : NAME`
- Or implementation: `module Name : NAME = struct ... end`

We will:

1. Build a quick and dirty on disk (key x value) storage.
2. Make the value type polymorph by requiring serializers..
3. Write an interface for it.
4. Then we will write a simple in memory storage.
5. And combine them into a cached disk storage.
6. We will make them available as three modules with the same interface.

A simple (key x value) database:

- One file for each key.
- String data stored as is.

```
1: let init dir =
2:   Unix.mkdir dir 0o750
3: let put n v dir =
4:   let fp = open_out (Filename.concat dir n) in
5:   output_string fp v ;
6:   close_out fp
7: let get n dir =
8:   let fn = Filename.concat dir n in
9:   if not (Sys.file_exists fn) then raise Not_found ;
10:  let fp = open_in fn in
11:  let len = in_channel_length fp in
12:  let buf = Bytes.create len in
13:  really_input fp buf 0 len ;
14:  close_in fp ;
15:  Bytes.to_string buf
```

We implement value polymorphism by providing converters:

```
1: let init dir =
2:   Unix.mkdir dir 0o750
3: let put to_string n v dir =
4:   let fp = open_out (Filename.concat dir n) in
5:   output_string fp (to_string v) ;
6:   close_out fp
7: let get of_string n dir =
8:   let fn = Filename.concat dir n in
9:   if not (Sys.file_exists fn) then raise Not_found ;
10:  let fp = open_in fn in
11:  let len = in_channel_length fp in
12:  let buf = Bytes.create len in
13:  really_input fp buf 0 len ;
14:  close_in fp ;
15:  of_string (Bytes.to_string buf)
```

We want to give it a proper abstract type interface:

```
1: type 'a table
2:
3: val init
4:   : ('a -> string) -> (string -> 'a) -> string -> 'a table
5: val put
6:   : string -> 'a -> 'a table -> unit
7: val get
8:   : string -> 'a table -> 'a
```


So we make a real 'a table type, in which we embed the converters:

```
1: type 'a table =  
2:   { to_string : 'a -> string;  
3:     of_string : string -> 'a ;  
4:     dir : string }  
5: let init to_string of_string dir =  
6:   Unix.mkdir dir 0o750 ;  
7:   { to_string ; of_string ; dir }  
8: let put n v { to_string ; dir } =  
9:   (* unchanged *)  
10: let get n { of_string ; dir } =  
11:   (* unchanged *)
```

Now let's write the memory storage, using OCaml hash tables.
We write it so that it fits in the same interface.

```
1: type 'a table =  
2:   { to_string : 'a -> string;  
3:     of_string : string -> 'a ;  
4:     table : (string, string) Hashtbl.t }  
5: let init to_string of_string _ =  
6:   let open Hashtbl in  
7:   let table : (string, string) t = create 100 in  
8:   { to_string ; of_string ; table }  
9: let put n v { to_string ; table } =  
10:  Hashtbl.replace table n (to_string v)  
11: let get n { of_string ; table } =  
12:  of_string (Hashtbl.find table n)
```

Notice how the database name is unused.
We'll correct that later on.

We group them inside a single compilation unit:

```
1: module In_memory_table = struct
2:   type 'a table = { (* ... *) } (* ... *)
3: end
4: module On_disk_table = struct
5:   type 'a table = { (* ... *) } (* ... *)
6: end
```

Whose interface is:

```
1: module In_memory_table : sig
2:   type 'a table (* ... *)
3: end
4: module On_disk_table : sig
5:   type 'a table (* ... *)
6: end
```

Of course, we want to make clear that the API is the same for both.
So we name it and the interface becomes:

```
1: module type TABLE = sig
2:   type 'a table
3:   val init
4:     : ('a -> string) -> (string -> 'a) -> string -> 'a table
5:   val put : string -> 'a -> 'a table -> unit
6:   val get : string -> 'a table -> 'a
7: end
8:
9: module In_memory_table : TABLE
10: module On_disk_table : TABLE
```

TABLE appears in the interface, it must appear in the implementation.

```
1: module type TABLE = sig (* ... *) end
2: module In_memory_table = struct (* ... *) end
3: module On_disk_table = struct (* ... *) end
```

Now we can add implementations of the API, for instance this cached one.

```
1: module Cached_table = struct
2:   type 'a table =
3:     'a On_disk_table.table * 'a In_memory_table.table
4:   let init to_string of_string dir =
5:     let ds = On_disk_table.init to_string of_string dir in
6:     let ms = In_memory_table.init to_string of_string dir in
7:     (ds, ms)
8:   let put n v (ds, ms) =
9:     On_disk_table.put n v ds ;
10:    In_memory_table.put n v ms
11:   let get n (ds, ms) =
12:     try
13:       In_memory_table.get n ms
14:     with Not_found ->
15:       let res = On_disk_table.get n ds in
16:       In_memory_table.put n res ms ;
17:       res
18: end
```

In the previous example, the database name was:

- used by the filesystem implementation as a relative directory ;
- dropped by the memory implementation.

Both cases are suboptimal.

Ideally, we would like to:

- Require all modules to provide a `param` type.
- While leaving implementations free to use whatever they need (here `unit` for in-memory, a `path` for on-disk).
- And making these types public.
(their values have to be forged externally to be passed to `init`).

That is what interface rewriting is for.

The abstract types of a named interface can be refined in two ways.

- Variant 1: `NAME with type t = def`
will build the same interface with `t` publicly specified to be `def`.
- Variant 2: `NAME with type t := def`
will remove the declaration of `t` and replace all its occurrences with `def`

Suppose we have the interface:

```
1: module type OF_STRING = sig
2:   type t
3:   val of_string : string -> t
4: end
```

We can rewrite it as follows:

```
1: module Float_of_string : OF_STRING with type t = float
```

Equivalent to writing

```
1: module Float_of_string : sig
2:   type t = float
3:   val of_string : string -> t
4: end
```


Or alternatively, we can make `t` disappear:

```
1 : module Float_of_string : OF_STRING with type t := float
```

Equivalent to writing

```
1 : module Float_of_string : sig
2 :   val of_string : string -> float
3 : end
```

Example: contextual table initialization

In our example, we could update `TABLE` to be:

```
1: module type TABLE = sig
2:   type 'a table
3:   type param
4:   val init
5:   : ('a -> string) -> (string -> 'a) -> param -> 'a table
6:   val put : string -> 'a -> 'a table -> unit
7:   val get : string -> 'a table -> 'a
8: end
```

Then the modules could be given refined signatures:

```
1: module In_memory_table
2:   : TABLE with type param := unit
3: module On_disk_table
4:   : TABLE with type param := string
```

An other module language primitive is `include`.

In signatures

- To type module level inheritance.
- With signature rewriting, to type module level traits.
- Syntax: `include NAME`
- With rewriting: `include NAME with type ...`

In implementation,

- To extend or patch existing modules.
- Syntax: `include Name`
- With signature: `include (Name : NAME)`
- With rewritten signature: `include (Name : NAME with type ...)`

We want to provide pre-instantiated cached tables for primitive types. Their interface is simpler, for instance `Int_table`'s should be:

```
1: module Int_table : sig  
2:   type t  
3:   val init : string -> t  
4:   val put : string -> int -> t -> unit  
5:   val get : string -> t -> int  
6: end
```

First, let's define the interface.

We define a generic interface:

```
1: module type TYPED_TABLE = sig
2:   type t
3:   type value
4:   val init : string -> t
5:   val put : string -> value -> t -> unit
6:   val get : string -> t -> value
7: end
```

And instantiate it by rewriting.

```
1: module Int_table
2:   : TYPED_TABLE with type value := int
3: module Float_table
4:   : TYPED_TABLE with type value := float
5: module String_table
6:   : TYPED_TABLE with type value := string
```

Then, we implement the modules, by patching the polymorphic implementation.

```
1: module Int_table = struct
2:   include In_memory_table
3:   type t = int table
4:   let init dir = init string_of_int int_of_string dir
5: end
6: module Float_table = struct
7:   include In_memory_table
8:   type t = float table
9:   let init dir = init string_of_float float_of_string dir
10: end
11: module String_table = struct
12:   include In_memory_table
13:   type t = string table
14:   let init dir = init (fun s -> s) (fun s -> s) dir
15: end
```

We want an alternative version of tables, with a `clear` operation. We extend the signature using `include`, and instantiate it:

```
1: module type CLEARABLE_TABLE = sig
2:   include TABLE
3:   val clear : 'a table -> unit
4: end
5:
6: module Clearable_in_memory_table
7:   : CLEARABLE_TABLE with type param := unit
8: module Clearable_on_disk_table
9:   : CLEARABLE_TABLE with type param := string
10: module Clearable_cached_table
11:   : CLEARABLE_TABLE with type param := string
```

And we extend the implementations using `include`:

```
1: module Clearable_in_memory_table = struct
2:   include In_memory_table
3:   let clear { table } = Hashtbl.clear table
4: end
5: module Clearable_on_disk_table = struct
6:   include On_disk_table
7:   let clear { dir } =
8:     Array.iter
9:       (fun f -> Unix.unlink (Filename.concat dir f))
10:      (Sys.readdir dir)
11: end
12: module Clearable_cached_table = struct
13:   include Cached_table
14:   let clear (ds, ms) =
15:     Clearable_in_memory_table.clear ms ;
16:     Clearable_on_disk_table.clear ds
17: end
```


Using Standard Functors

Functor Application
Set and Map

Functor Application

What are functors ?

- Parametric modules.
- Module level functions.

The code of a functor:

- Supposes that some module of a given signature exists.
- Is linked to such a module by a **functor application** at runtime.

Syntax:

- Functor signatures:
`module Name : functor (Arg : ARG) -> RESULT.`
- Functor application:
`module Instance = Name (Arg).`

The standard library defines:

- `Set`, functional sets of elements of a given type ;
- `Map` functional polymorphic maps of keys of a given type.

For constructing maps, the `map.mli` exports the functor:

```
1 : module Make (Ord : OrderedType) : S with type key = Ord.t
```

- `OrderedType` is the signature that the parameter must respect.
It defines an abstract type `t`, the type of keys.
- `S` is the signature of the result.
It is linked to the parameter by the `with type` syntax.

In more details, the `OrderedType` signature is:

```
1 : module type OrderedType = sig
2 :   type t
3 :   val compare : t -> t -> int
4 : end
```

And an extract of `S`:

```
1 : module type S = sig
2 :   type key
3 :   type (+ 'a) t
4 :   val empty: 'a t
5 :   val is_empty: 'a t -> bool
6 :   (* ... *)
7 : end
```

Let's instantiate `Map` to build maps of strings.

```
1: module StringOrderedType = struct  
2:   type t = string  
3:   let compare = Pervasives.compare  
4: end  
5: module StringMap = Map.Make (String)
```

`StringMap` is then usable as any other module.

More concisely:

```
1 : module StringMap = Map.Make (struct
2 :   type t = string
3 :   let compare = Pervasives.compare
4 : end)
```

Actually, the `String` module already defines `t` and `compare`:

```
1 : module StringMap = Map.Make (String)
```

If a map is only needed for a local computation, one can write:

```
1 : let module StringMap = Map.Make (String) in (* expr *)
```

Outline

Toplevel Modules
Abstraction
The module Language
Using Standard Functors