

Memory management

OCaml **PRO** – Grégoire Henry

March 9-13 2015

Outline

Uniform data representation
Garbage collection

Uniform data representation

Memory graph traversal

Expected properties for a Garbage Collector (GC):

- deallocate **unused** heap values as soon as possible;
- do not deallocate values **required** by further program execution;
- a fast allocation mechanism;
- do not block the program **too long** while collecting.

The GC **traverses the memory graph** to identify:

alive values that may be accessed from the roots (globals, stack) and potentially used later.

dead values : all others values, a subset of **unused** values.

A GC is:

precise if the memory graph can be traversed without ambiguity.

conservative (ambiguous), if it has to over-estimate the set of living value.

⇒ the GC of OCaml is precise.


Uniform data representation

- every value is represented with a single word;
- some immediate values fit in a word (int, char);
- bigger values are heap-allocated (a.k.a. **boxed values**).

This allows to compile parametric polymorphism with code sharing
There is no specialisation by default.

An OCaml value is either:

- a word-aligned pointer (Least significant bit: 0), or
- an even integer (Least significant bit: 1).




All pointers should target the OCaml heap.

Out of heap pointers are not valid immediate values, and should be properly boxed (since OCaml 4.02).

All pointers should target a **block**, starting with a single word header containing:

- a tag (8 bits);
- GC bits, a.k.a. a color (2 bits);
- the block size in words (22 or 54 bits).



Maximum size of a block: 2^{22} words \approx 16MB (32 bits), 2^{54} words \approx a lot (64 bits)

The block header tag is used by the GC to discriminate:

- Scannable block: every word in the block is an OCaml value.
 - 0-245** Generic block (sum type)
 - 246** Non-evaluated lazy value
 - 247** Closure
 - 248** Object
 - 249** Infix (mutually recursive closure)
 - 250** Forward (evaluated lazy value, internal usage)
- Non-scannable block: they must not contain pointers to the OCaml heap.
 - 251** Abstract value (may contains out of heap pointers)
 - 252** String
 - 253** Double
 - 254** Double array
 - 255** Custom block (may contains out of heap pointers)

This is enough information to precisely traverse
the memory graph from its roots: global value, stack frame, ...

int an immediate value

char an immediate value (8 significant bits)

bool an immediate value (only 1 significant bit!)

float^a a boxed double (3 words)


int32^a, **int64**^a, **nativeint**^a boxed values (custom block: 4, 6 or 4/6 words)

array a block: tag 0, $n+1$ words (when `Array.length = n`)

float array^a a block: $2*n+1$ words (when `Array.length = n`)

tuples/record types a block: tag 0, $n+1$ words (when n is the number of fields)

```
1: type t = { a: int; b: float; }
2: let x = { a = 1; b = 1.0; }
3: type t2 = { a: bool; b: float; }
4: let x2 = { a = true; b = 1.0; }
```



^aCompiler contains some optimisation to locally avoid boxing/unboxing.

sum types


constant constructor an immediate value i , when the constructor is the i -th constant constructor in the type declaration;

non-constant constructor a block

- tag i , when the constructor is the i -th non-constant constructor in the type declaration;
- size = the number of parameter.

1 : **type** t = A **of** t * t | B | C **of** (t * t) | D | E

2 : **let** x = A (B, C (D, E))



string/bytes


- a block of tag 252.
- size, when `String.length s = n`:

$$1 + \frac{n+1}{8}$$

- the last byte of the block contains:

$$8 * \text{size} - (n + 1)$$


```
1 : let x = "123456"      (* String.length x = 6 *)
2 : let y = "1234567"    (* String.length y = 7 *)
3 : let z = "123456\000" (* String.length z = 7 *)
```



function (intuitively)

- a block of tag 247
- first field: code pointer
- second field: arity
- remaining fields: the captured environment (no globals)

```
1: let x = 3
2: let make y =
3:   let cpt = ref 0 in
4:   fun () -> cpt := !cpt + x + y; !cpt
5: let f = make 5
```



partial application a closure containing the original closure and the applied arguments

mutually recursive function a shared closure for all the functions

module a block

functor a closure

How many heap allocated words?

```
1 : let x = [1;2;3]
2 : let y = [1.;2.;3.]
3 : let z = ["1"; "2"; "3"]
4 :
5 : let a = [|1;2;3|]
6 : let b = [|1.;2.;3.|]
7 : let c = [| "1"; "2"; "3" |]
8 :
9 : let t = (1, 2., '3')
10 :
11 : let rec 1 = 1 :: 2 :: 3 :: 1
```

Value introspection: the non-documented module Obj.

```
1 : (** Not for the casual user. *)
2 :
3 : type t
4 :
5 : val repr : 'a -> t
6 :
7 : val is_block : t -> bool
8 : val is_int : t -> bool
9 : val tag : t -> int
10 : val size : t -> int
11 : val field : t -> int -> t
12 : val double_field : t -> int -> float
13 :
14 : val new_block : int -> int -> t
15 : val set_field : t -> int -> t -> unit
16 : val set_double_field : t -> int -> float -> unit
```

Garbage collection

OCaml has a generational GC:


- First generation (a.k.a. **minor heap**): stop-and-copy
- Second generation (a.k.a. **major heap**):
 - incremental mark-and-sweep (most of the time)
 - stop-and-copy (when a compaction is required)

The GC is tailored for functional programming and symbolic processing :

- a lot of short-lived value;
- a lot of small values;
- few mutations.

Minor heap

Immutable blocks smaller than a constant (currently 256) are allocated on the minor heap. Allocating in the minor heap is as fast as allocation in the stack with C++!



- Allocation is simply incrementing current.
- When the minor heap is full, live blocks (■)
- Dead blocks (■) “disappears” when current is reverted to minor_start.

Mutable blocks and huge blocks are allocated in the major heap.

- Allocation by traversing the (lonely) free-list. Two strategies:
 - Next-fit** start the traversal from the last allocated value (fast allocation);
 - First-fit** start the traversal from the beginning of the list (prevents unbounded fragmentation).
- Three phases:
 - Idle** Do nothing!
 - Marking** Traverse the graph to mark live blocks.
 - Sweeping** Linearly scan the heap to free non-live blocks.
- Incremental: after a minor collection, execute a small part of the current phase.
- Compaction: when the fragmentation is too important (may be disabled).

How to handle pointer between generations ?

- Pointer from major heap to minor heap are required to detect living blocks in the minor heap.
 - traversing the whole major heap for a minor collection would be too costly;
 - they may only appear while mutating an “old” object from the major heap;

The GC keep a list of such pointers: at every mutation in the major heap should test the “age” of the written value (still a little bit costly).

The reverse list, is also the required to update the pointer when a block is copied from the minor heap to the major heap.

- Pointer from minor heap to major heap are “removed” while copying blocks from the minor heap to the major heap.

Tweaking & introspection

The GC has some parameters that may be tuned:

minor_heap_size current default: 256k

major_heap_increment in percent, default: 15%

space_overhead major GC eagerness

max_overhead ratio waste/live that trigger compaction, default: 500%

stack_limit default: 1M (bytecode only)

allocation_policy default: next fit

verbose

Modified by settings the OCAMLRUNPARAM environment variable, or dynamically by calling the function `Gc.set`.

```
1 : Gc.set
2 :   { Gc.get () with
3 :     Gc.minor_heap_size = 1024 * 1024;
4 :     Gc.max_overhead = 1_000_000; (* disable compaction *)
5 :     Gc.allocation_policy = 1; (* First-fit *) }
```